


● Full-S1 ■ K1 FP32 ▲ Q8(K1) simulated



Secondary extension. Post-training per-output-channel symmetric int8 weights; 95.87% of parameters stored as int8; CPU torch.ao dynamic backend; no true CUDA int8 path. The Q8 non-inferiority result applies to the weight-only simulated representation, not directly to the deployed cpu_dynamic artifact.